


**VISVESVARAYATECHNOLOGICALUNIVERSITY
JNANASANGAMA, BELGAUM-590014**

ADDITIONAL ACTIVITY-1


FOUNDATIONS OF DATA SCIENCE

INTRODUCTION TO R PROGRAMMING



Team member:

Yashaswini MB	1SI19CS139
Yashaswini GR	1SI19CS143
Shruthi N	1SI19CS145



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

SIDDAGANGA INSTITUTE OF TECHNOLOGY, TUMKUR-572103

(An Autonomous Institute, Affiliated to Visvesvaraya Technological University, Belgaum, Recognized by

AICTE and Accredited by NBA, New Delhi)

2021-2022

ACTIVITY 1

INTRODUCTION

R programming language is an open source programming language that has been widely used across the world. It is the language that can also help businesses analyze vast amounts of information quickly and effectively.

2 Million Plus Users across the world. And the number of users is steadily growing. R has set its foothold in the Data Science industry and owing to its massive repository, it has become the most sought-after programming language in the world.

WHAT IS R PROGRAMMING?

R is an open source programming language used for statistical computing. It is one of the most popular programming languages today. R was inspired by S+, it is similar to the S programming language.

FEATURES OF R

- **Open-source**

R is an open-source programming language. This means that it is free of cost and requires no license.

- **Comprehensive Language**

R is a comprehensive programming language, meaning that it provides services for statistical modeling as well as for software development. R is the primary language for Data Science as well as for developing web applications.

- **Provides a Wide Array of Packages**

R is most widely used because of its wide availability of libraries. R has CRAN, which is a repository holding more than 10,000 packages.

- **Possesses a Number of Graphical Libraries**

The most important feature of R that sets it apart from other programming languages of Data Science is its massive collection of graphical libraries like ggplot2, plotly, etc. that are capable of making aesthetic and quality visualizations.

- **Performs Fast Calculations**

Through R, you can perform a wide variety of complex operations on vectors, arrays, data frames and other data objects of varying sizes. Furthermore, all these operations operate at a lightning speed.

CONCEPTS IN R

Variables

Variables can be created using the "<-" (assignment) operator.

Variables are case-sensitive. Var1 and var1 are considered different variables.

The class function can be used on variables to determine the data type of the variable which is classified in three major types – character, numeric and logical.

```
myString <- "Hello, World!"
```

If Else Statement

An if statement can be followed by an optional else statement which executes when the boolean expression is false.

Syntax

```
if(boolean_expression)
{ // statement(s) will execute if the boolean expression is true. }
else
{ // statement(s) will execute if the boolean expression is false. }
```

If the Boolean expression evaluates to be true, then the if block of code will be executed, otherwise else block of code will be executed.

While Loop

The While loop executes the same code again and again until a stop condition is met

Syntax

```
while (test_expression) {
  statements
}
```

Function Definition

A function is a set of statements organized together to perform a specific task.

An R function is created by using the keyword function.

syntax

```
function_name <- function(arg_1, arg_2, ...){
  Function body
}
```

Paste function

Paste function in R is used to concatenate Vectors by converting them into character. It also used to concatenate the two string values by separating with delimiters specified by keyword sep. Example

```
paste('one',2,'three',4,'five',sep=',')
[1] "one,2 ,three,4,five"
```

Floor function

The R floor method is one of the R Math functions, which is to return the largest integer value. That is not greater than (less than) or equal to a specific number or an expression.

Example

```
print(floor(2.6))
[1] 2
```

Caret symbol

It is the exponent operator. calculate a base number raised to the power of exponent number.

Example

```
print(2^3)
[1] 8
```

PROGRAM TO CHECK WHETHER GIVEN NUMBER IS ARMSTRONG OR NOT

```
isArmstrong <- function(x){
  sum <- 0
  y <- x
  while(y > 0){
    digit <- y %% 10
    sum <- sum + (digit^3)
    y <- floor(y/10)
  }
  if(x == sum){
    print(paste(x,' is an Armstrong Number '))
  } else {
    print(paste(x,'is not an Armstrong Number'))
  }
}
# take input from the user
num = as.integer(readline(prompt="Enter a number: "))
isArmstrong (num)
```

Explanation:

To check whether the number is armstrong or not, we have defined a function isArmstrong with x as the argument here x is the number which we have to check if it is Armstrong or not.

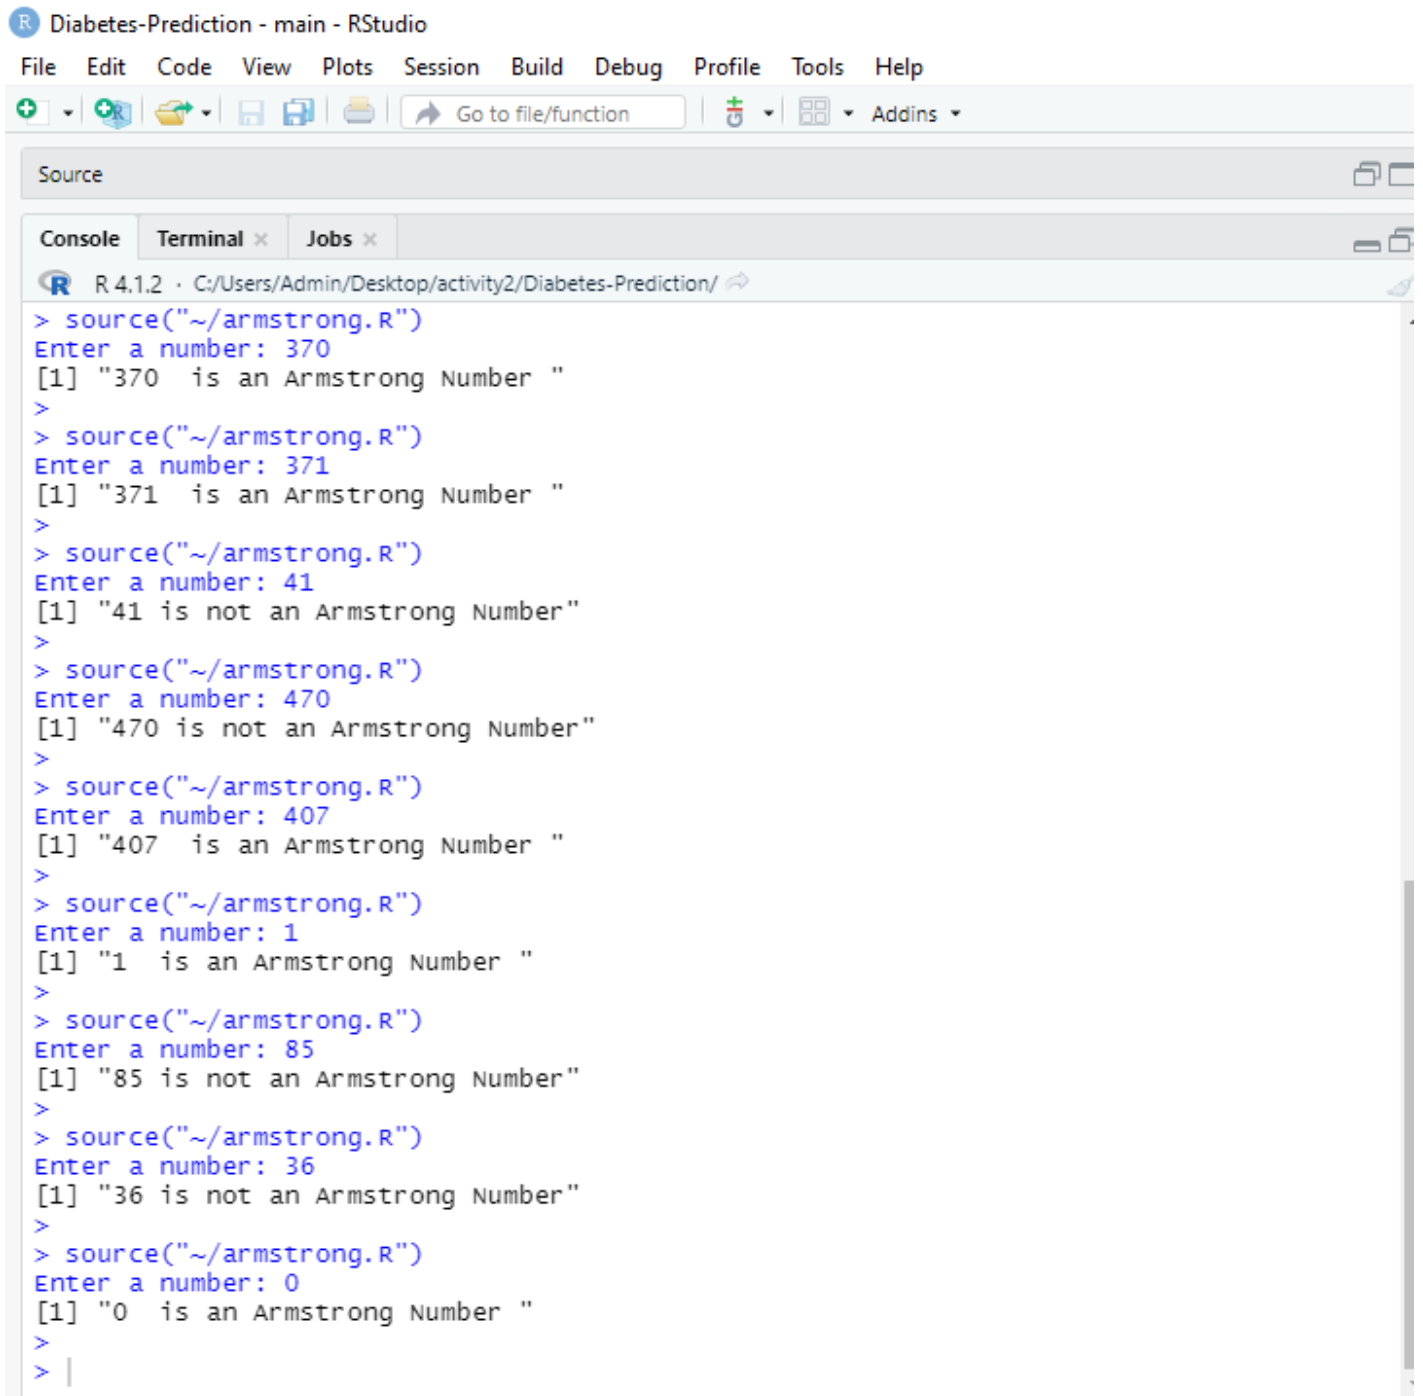
We need to calculate the sum of the cube of each digit. So, we initialize the sum to 0 and obtain each digit number by using the **modulus operator** %%.

The remainder of a number when it is divided by 10 is the last digit of that number. We take the cubes using **exponent operator**.

Finally, we compare the sum with the original number and conclude that it is Armstrong number if they are equal.

To read the input from user the **prompt** argument is chosen to display an appropriate message for the user. Use **readline()** function to take input from the user (terminal). And we convert the input, which is a character vector into integer using the function **as.integer()**.

OUTPUT



The screenshot shows the RStudio interface with the following components:

- Title Bar:** R Diabetes-Prediction - main - RStudio
- Menu Bar:** File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help
- Toolbar:** Includes icons for opening files, saving, printing, and a search bar labeled "Go to file/function".
- Source Panel:** Shows the file path: C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/
- Console Panel:** Displays the following R session output:

```
> source("~/armstrong.R")
Enter a number: 370
[1] "370 is an Armstrong Number "
>
> source("~/armstrong.R")
Enter a number: 371
[1] "371 is an Armstrong Number "
>
> source("~/armstrong.R")
Enter a number: 41
[1] "41 is not an Armstrong Number"
>
> source("~/armstrong.R")
Enter a number: 470
[1] "470 is not an Armstrong Number"
>
> source("~/armstrong.R")
Enter a number: 407
[1] "407 is an Armstrong Number "
>
> source("~/armstrong.R")
Enter a number: 1
[1] "1 is an Armstrong Number "
>
> source("~/armstrong.R")
Enter a number: 85
[1] "85 is not an Armstrong Number"
>
> source("~/armstrong.R")
Enter a number: 36
[1] "36 is not an Armstrong Number"
>
> source("~/armstrong.R")
Enter a number: 0
[1] "0 is an Armstrong Number "
>
> |
```

ACTIVITY 2

INTRODUCTION

Diabetes:

Diabetes is a standout amongst the most well-known non-transmittable diseases in the world. It is assessed to be the seventh leading cause for death. Diabetes causes a large number of deaths each year and a large number of people living with the disease do not realize their health condition early enough.

According to a study by the World Health Organization (WHO), this number will have raised to 552 million by 2030, denote that one in 10 grownups will have diabetes by 2030 if no serious act is taken. In 2014, the worldwide frequency of diabetes was projected to be 9 % among adults aged 18+ years

Diabetes is a disease caused due to the increase level of blood glucose. Diabetes is a chronic disease with the potential to cause a worldwide health care crisis.

Various traditional methods, based on physical and chemical tests, are available for diagnosing diabetes. However, early prediction of diabetes is quite challenging task for medical practitioners due to complex interdependence on various factors as diabetes affects human organs such as kidney, eye, heart, nerves, foot etc. Data science methods have the potential to benefit other scientific fields by shedding new light on common questions.

One such task is to help make predictions on medical data. Machine learning is an emerging scientific field in data science dealing with the ways in which machines learn from experience.

OBSERVATION

The number of people with diabetes has risen from 108 million in 1980 to 422 million in 2014, with the global prevalence of diabetes among adults over 18 years of age rising from 4.7% in 1980 to 8.5% in 2014. By 2040, 642 million adults (1 in 10 adults) are expected to have diabetes. Also, 46.5% of those with diabetes have not been diagnosed. In order to reduce the number of deaths attributable to diabetes, it is essential that methods and techniques that will aid in early diagnosis of diabetes be devised, because a large number of deaths in diabetic patients are due to late diagnosis.

PROBLEM STATEMENT

The improved model for early diabetes prediction by integrating Logistic Regression and Decision Tree techniques.

APPROACHES

- **Logistic Regression**

Logistic regression is a classification model in machine learning, extensively used in clinical analysis. It uses probabilistic estimations which helps in understanding the relationship between the dependent variable and one or more independent variables. Diabetes, being one of the most common diseases around the world, when detected early, may prevent the progression of the disease and avoid other complications. In this work, we design a prediction model, that predicts whether a patient has diabetes, based on certain diagnostic measurements included in the dataset, and explore various techniques to boost the performance and accuracy.

- **Decision Tree**

Decision Trees are a non-parametric supervised learning method used for classification and regression. The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features. A tree can be seen as a piecewise constant approximation.

DATASET

The data set used for the purpose of this study is Pima Indians Diabetes Database of National Institute of Diabetes and Digestive and Kidney Diseases. This diabetes database is a collection of medical diagnostic reports of 768 examples from a population living near Phoenix, Arizona, USA. The samples consist of examples with 8 attribute values and one of the two possible outcomes, namely whether the patient is tested positive for diabetes (indicated by output one) or not (indicated by zero).

Attribute Information:

1. Number of times pregnant
2. Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. Diastolic blood pressure (mm Hg)
4. Triceps skin fold thickness (mm)
5. 2-Hour serum insulin (μ U/ml)
6. Body mass index (weight in kg/(height in m)²)
7. Diabetes pedigree function
8. Age (years)
9. Class variable (0 or 1)

DESCRIPTIVE STATISTICS

IMPORTING LIBRARIES:

Libraries for Box Plot and Histogram:

```
library(ggplot2)
```

```
library(dplyr)
```

```
library(gridExtra)
```

```
library(corrplot)
```

Libraries for Class Distribution:

```
library(car)
```

```
library(caret)
```

```
library(class)
```

```
library(dplyr)
```

```
library(ggplot2)
```

```
library(tidyr)
```

```
library(tidyverse)
```

```
library(performance)
```


DATA SUMMARY

R Diabetes-Prediction - main - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

+ - R Go to file/function Addins

Source

Console

Terminal x

Jobs x

R 4.1.2 C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/

> summary(diabetes)

Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
Min. : 0.000	Min. : 0.0	Min. : 0.00	Min. : 0.00	Min. : 0.0
1st Qu.: 1.000	1st Qu.: 99.0	1st Qu.: 62.00	1st Qu.: 0.00	1st Qu.: 0.0
Median : 3.000	Median :117.0	Median : 72.00	Median :23.00	Median : 30.5
Mean : 3.845	Mean :120.9	Mean : 69.11	Mean :20.54	Mean : 79.8
3rd Qu.: 6.000	3rd Qu.:140.2	3rd Qu.: 80.00	3rd Qu.:32.00	3rd Qu.:127.2
Max. :17.000	Max. :199.0	Max. :122.00	Max. :99.00	Max. :846.0

BMI	DiabetesPedigreeFunction	Age	Outcome
Min. : 0.00	Min. :0.0780	Min. :21.00	Min. :0.000
1st Qu.:27.30	1st Qu.:0.2437	1st Qu.:24.00	1st Qu.:0.000
Median :32.00	Median :0.3725	Median :29.00	Median :0.000
Mean :31.99	Mean :0.4719	Mean :33.24	Mean :0.349
3rd Qu.:36.60	3rd Qu.:0.6262	3rd Qu.:41.00	3rd Qu.:1.000
Max. :67.10	Max. :2.4200	Max. :81.00	Max. :1.000

> describe(diabetes)

	vars	n	mean	sd	median	trimmed	mad	min	max	range
Pregnancies	1	768	3.85	3.37	3.00	3.46	2.97	0.00	17.00	17.00
Glucose	2	768	120.89	31.97	117.00	119.38	29.65	0.00	199.00	199.00
BloodPressure	3	768	69.11	19.36	72.00	71.36	11.86	0.00	122.00	122.00
SkinThickness	4	768	20.54	15.95	23.00	19.94	17.79	0.00	99.00	99.00
Insulin	5	768	79.80	115.24	30.50	56.75	45.22	0.00	846.00	846.00
BMI	6	768	31.99	7.88	32.00	31.96	6.82	0.00	67.10	67.10
DiabetesPedigreeFunction	7	768	0.47	0.33	0.37	0.42	0.25	0.08	2.42	2.34
Age	8	768	33.24	11.76	29.00	31.54	10.38	21.00	81.00	60.00
Outcome	9	768	0.35	0.48	0.00	0.31	0.00	0.00	1.00	1.00

	skew	kurtosis	se
Pregnancies	0.90	0.14	0.12
Glucose	0.17	0.62	1.15
BloodPressure	-1.84	5.12	0.70
SkinThickness	0.11	-0.53	0.58
Insulin	2.26	7.13	4.16
BMI	-0.43	3.24	0.28
DiabetesPedigreeFunction	1.91	5.53	0.01
Age	1.13	0.62	0.42
Outcome	0.63	-1.60	0.02

> |

Source

Console Terminal x Jobs x

R 4.1.2 · C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/

4 0.167 21 0

5 2.288 33 1

6 0.201 30 0

> cat("Number of missing value:", sum(is.na(diabetes)), "\n")

Number of missing value: 0

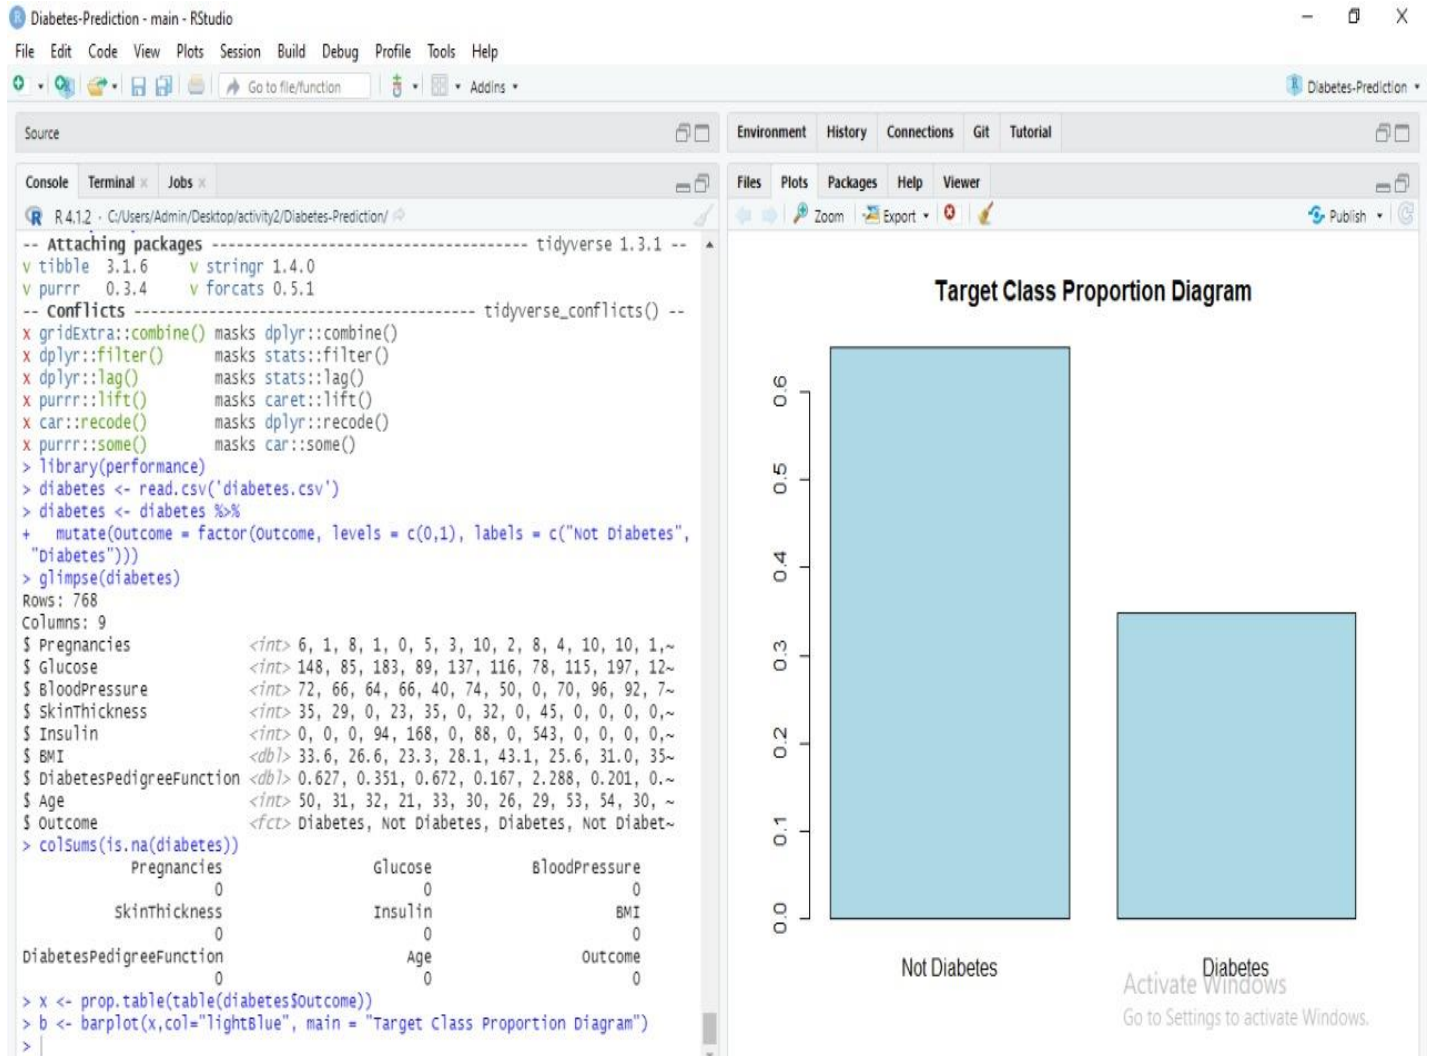
> summary(diabetes)

Pregnancies	Glucose	BloodPressure	SkinThickness
Min. : 0.000	Min. : 0.0	Min. : 0.00	Min. : 0.00
1st Qu.: 1.000	1st Qu.: 99.0	1st Qu.: 62.00	1st Qu.: 0.00
Median : 3.000	Median : 117.0	Median : 72.00	Median : 23.00
Mean : 3.845	Mean : 120.9	Mean : 69.11	Mean : 20.54
3rd Qu.: 6.000	3rd Qu.: 140.2	3rd Qu.: 80.00	3rd Qu.: 32.00
Max. : 17.000	Max. : 199.0	Max. : 122.00	Max. : 99.00

Insulin	BMI	DiabetesPedigreeFunction	Age
Min. : 0.0	Min. : 0.00	Min. : 0.0780	Min. : 21.00
1st Qu.: 0.0	1st Qu.: 27.30	1st Qu.: 0.2437	1st Qu.: 24.00
Median : 30.5	Median : 32.00	Median : 0.3725	Median : 29.00
Mean : 79.8	Mean : 31.99	Mean : 0.4719	Mean : 33.24
3rd Qu.: 127.2	3rd Qu.: 36.60	3rd Qu.: 0.6262	3rd Qu.: 41.00
Max. : 846.0	Max. : 67.10	Max. : 2.4200	Max. : 81.00

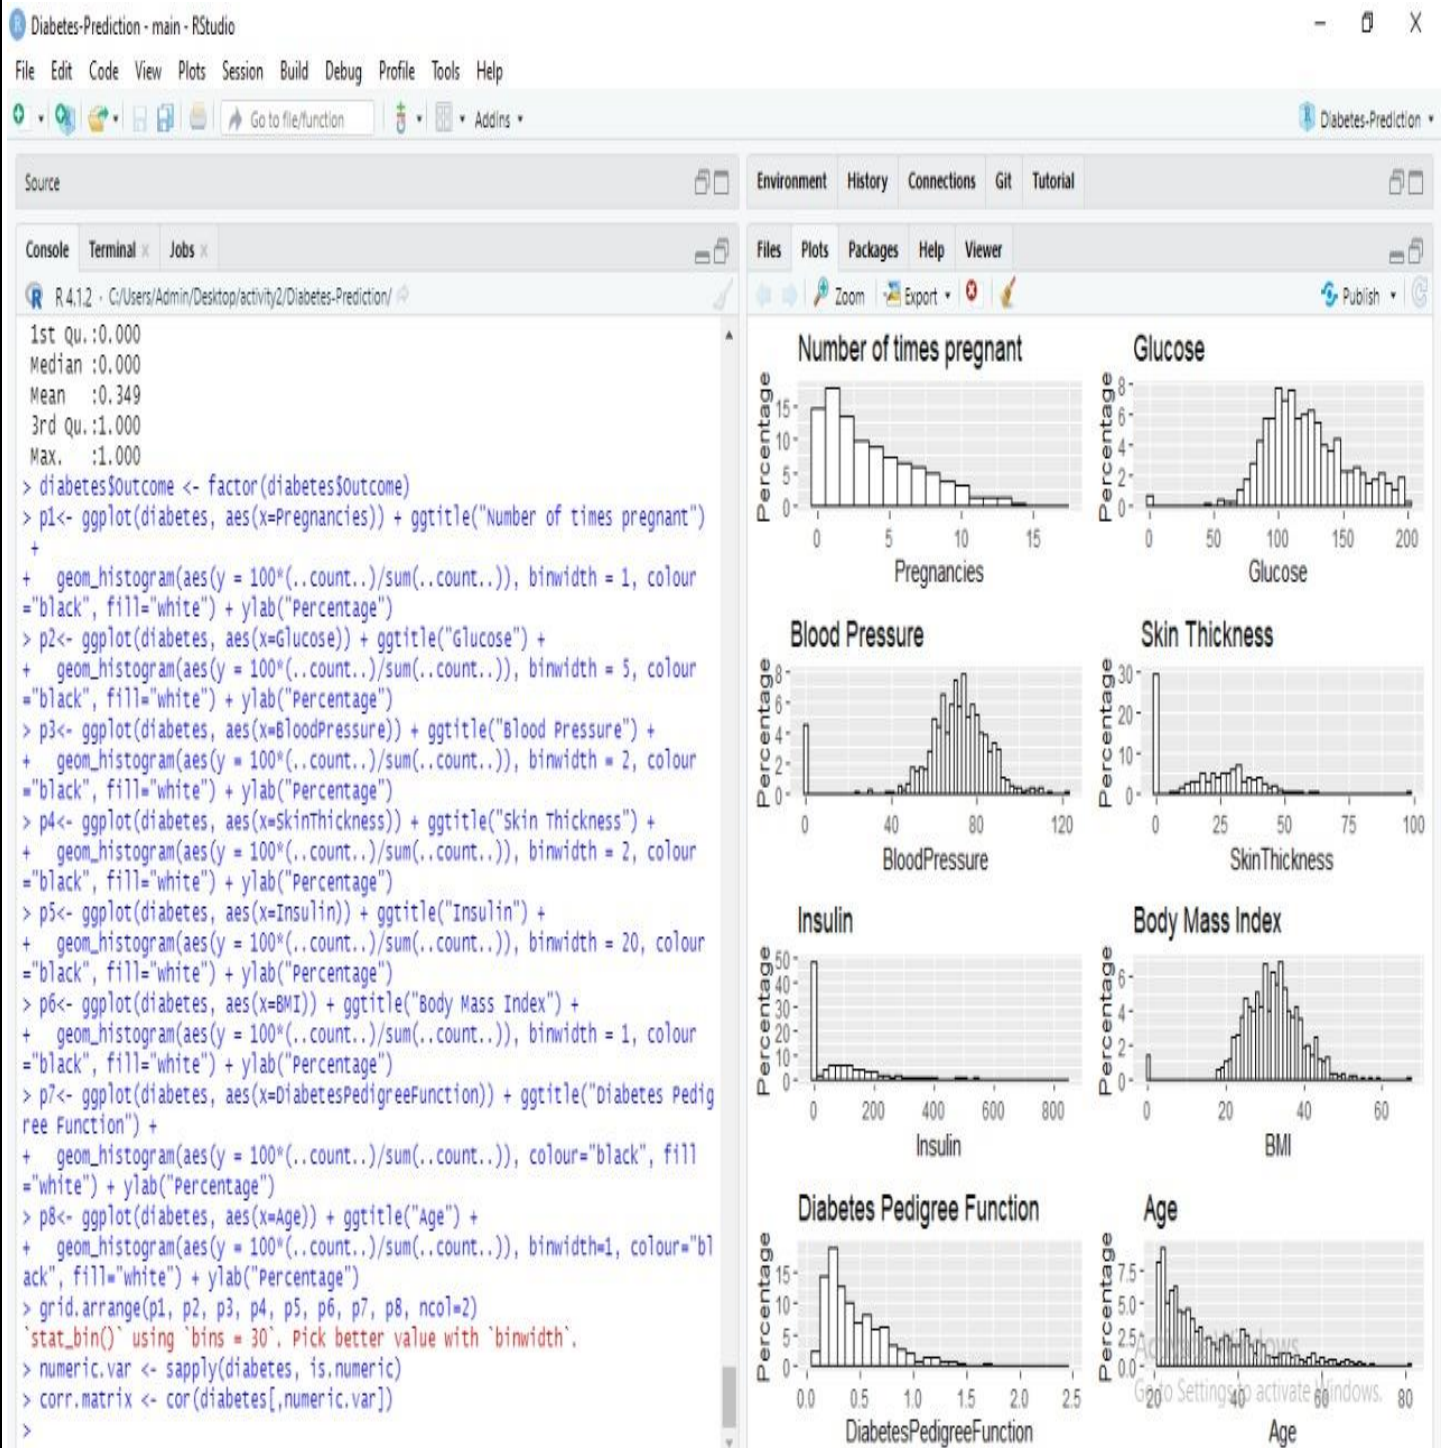
Outcome
Min. : 0.000
1st Qu.: 0.000
Median : 0.000
Mean : 0.349
3rd Qu.: 1.000
Max. : 1.000

CLASS DISTRIBUTION

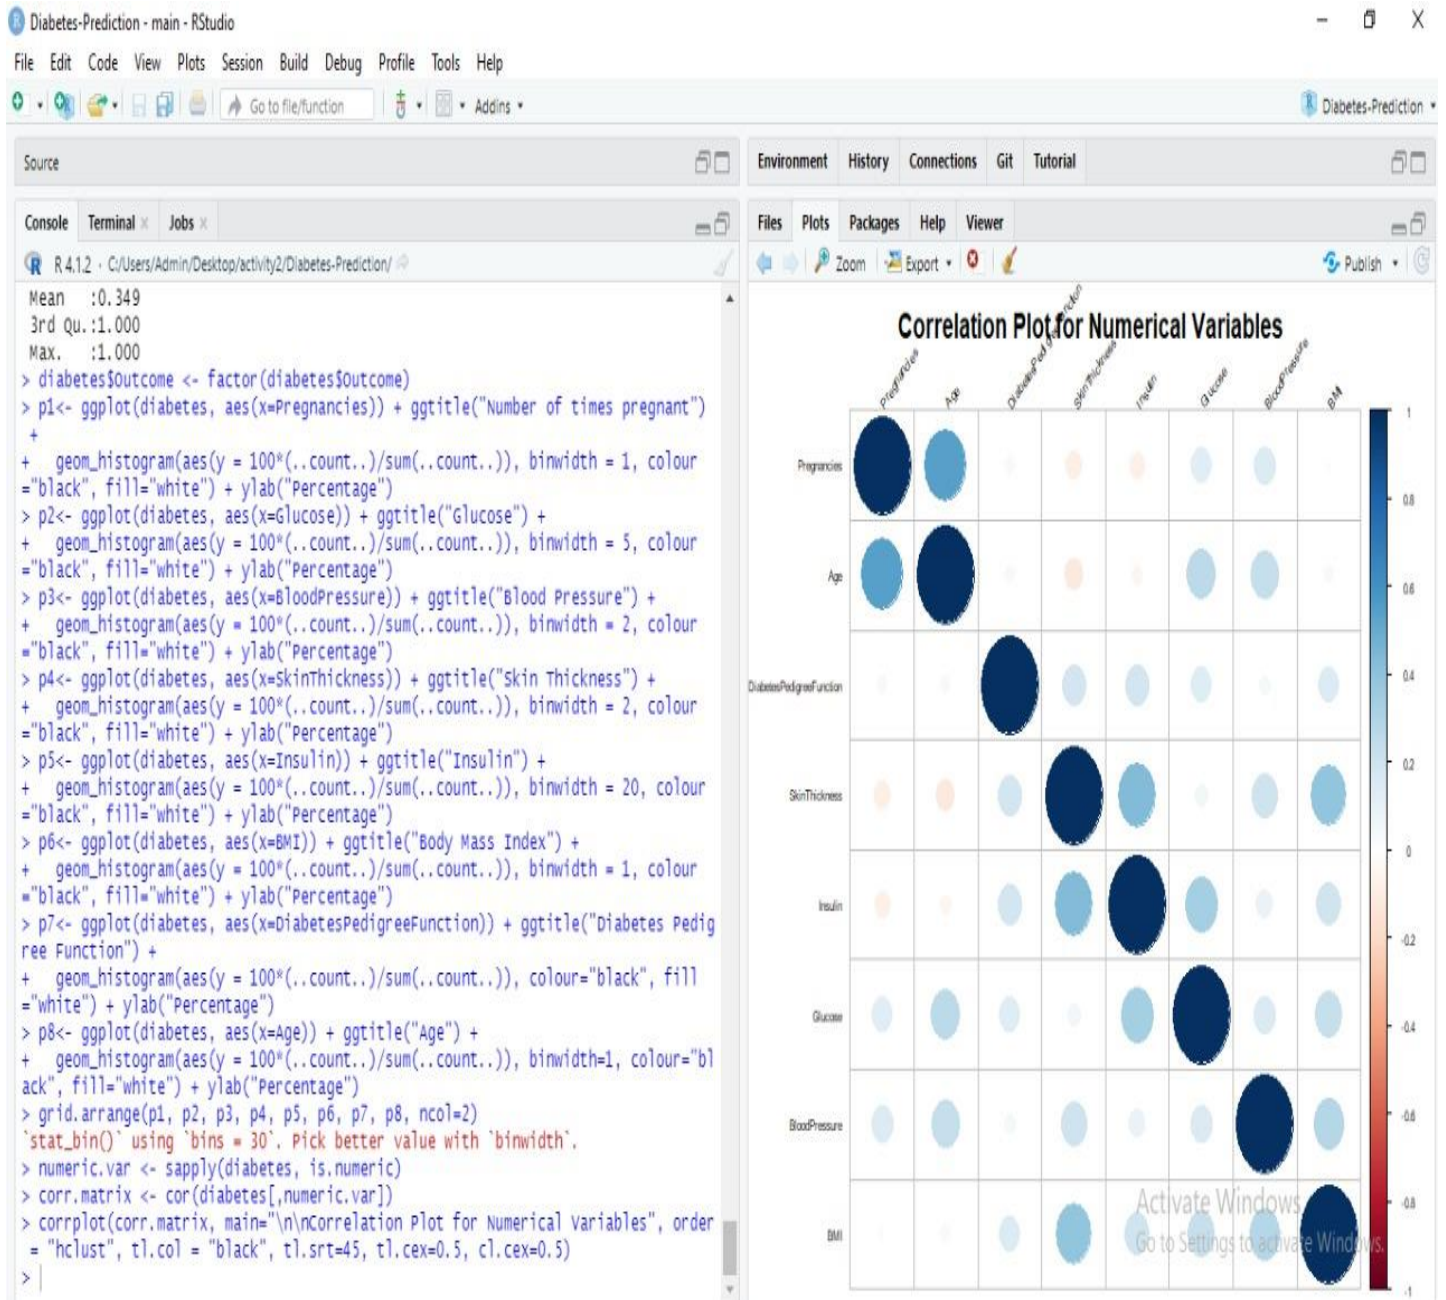


DATA VISUALISATION

Histogram Plot:



Correlation Plot:



Box and Whisker Plot:

Diabetes-Prediction - main - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

Diabetes-Prediction

Source Environment History Connections Git Tutorial

Console Terminal Jobs Files Plots Packages Help Viewer

R 4.1.2 C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/ Zoom Export Publish


```
+ geom_histogram(aes(y = 100*(..count../sum(..count..)), binwidth=1, colour="black", fill="white") + ylab("Percentage")
> grid.arrange(p1, p2, p3, p4, p5, p6, p7, p8, ncol=2)
'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
> numeric.var <- sapply(diabetes, is.numeric)
> corr.matrix <- cor(diabetes[,numeric.var])
> corplot(corr.matrix, main="\nCorrelation Plot for Numerical Variables", order = "hclust", tl.col = "black", tl.srt=45, tl.cex=0.5, cl.cex=0.5)
> attach(diabetes)
```

The following objects are masked from diabetes (pos = 3):

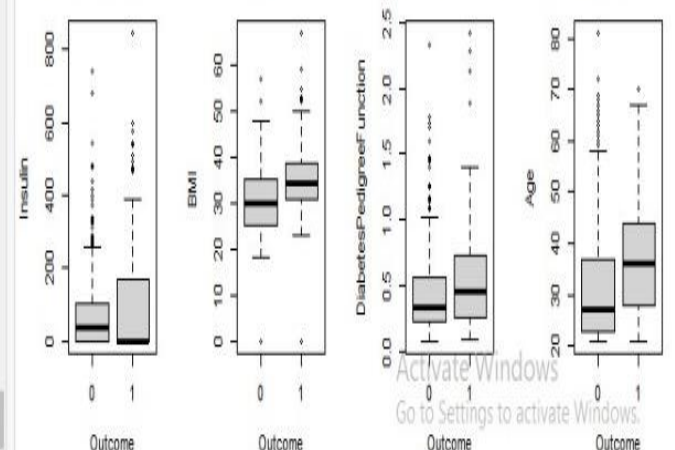
Age, BloodPressure, BMI, DiabetesPedigreeFunction, Glucose, Insulin, Outcome, Pregnancies, SkinThickness

```
> par(mfrow=c(2,4))
> boxplot(Pregnancies~Outcome, main="No. of Pregnancies vs. Diabetes",
+         xlab="Outcome", ylab="Pregnancies")
> boxplot(Glucose~Outcome, main="Glucose vs. Diabetes",
+         xlab="Outcome", ylab="Glucose")
> boxplot(BloodPressure~Outcome, main="Blood Pressure vs. Diabetes",
+         xlab="Outcome", ylab="Blood Pressure")
> boxplot(SkinThickness~Outcome, main="Skin Thickness vs. Diabetes",
+         xlab="Outcome", ylab="Skin Thickness")
> boxplot(Insulin~Outcome, main="Insulin vs. Diabetes",
+         xlab="Outcome", ylab="Insulin")
> boxplot(BMI~Outcome, main="BMI vs. Diabetes",
+         xlab="Outcome", ylab="BMI")
> boxplot(DiabetesPedigreeFunction~Outcome, main="Diabetes Pedigree Function vs. Diabetes",
+         xlab="Outcome", ylab="DiabetesPedigreeFunction")
> boxplot(Age~Outcome, main="Age vs. Diabetes",
+         xlab="Outcome", ylab="Age")
>
>
>
>
>
>
```

No. of Pregnancies vs. Diabetes Glucose vs. Diabetes Blood Pressure vs. Diabetes Skin Thickness vs. Diabetes



Insulin vs. Diabetes BMI vs. Diabetes Diabetes Pedigree Function vs. Diabetes Age vs. Diabetes



Activate Windows
Go to Settings to activate Windows.

PREDICTION ON DATA

Logistic regression:

we use a training data set containing a random sample of 70% of the observation to perform a Logistic Regression with “Diabetes” as the response and the remains variables as predictors.

Algorithm of Logistic Regression in R:

```
x <- cbind(x_train,y_train)

# Train the model using the training sets and check score

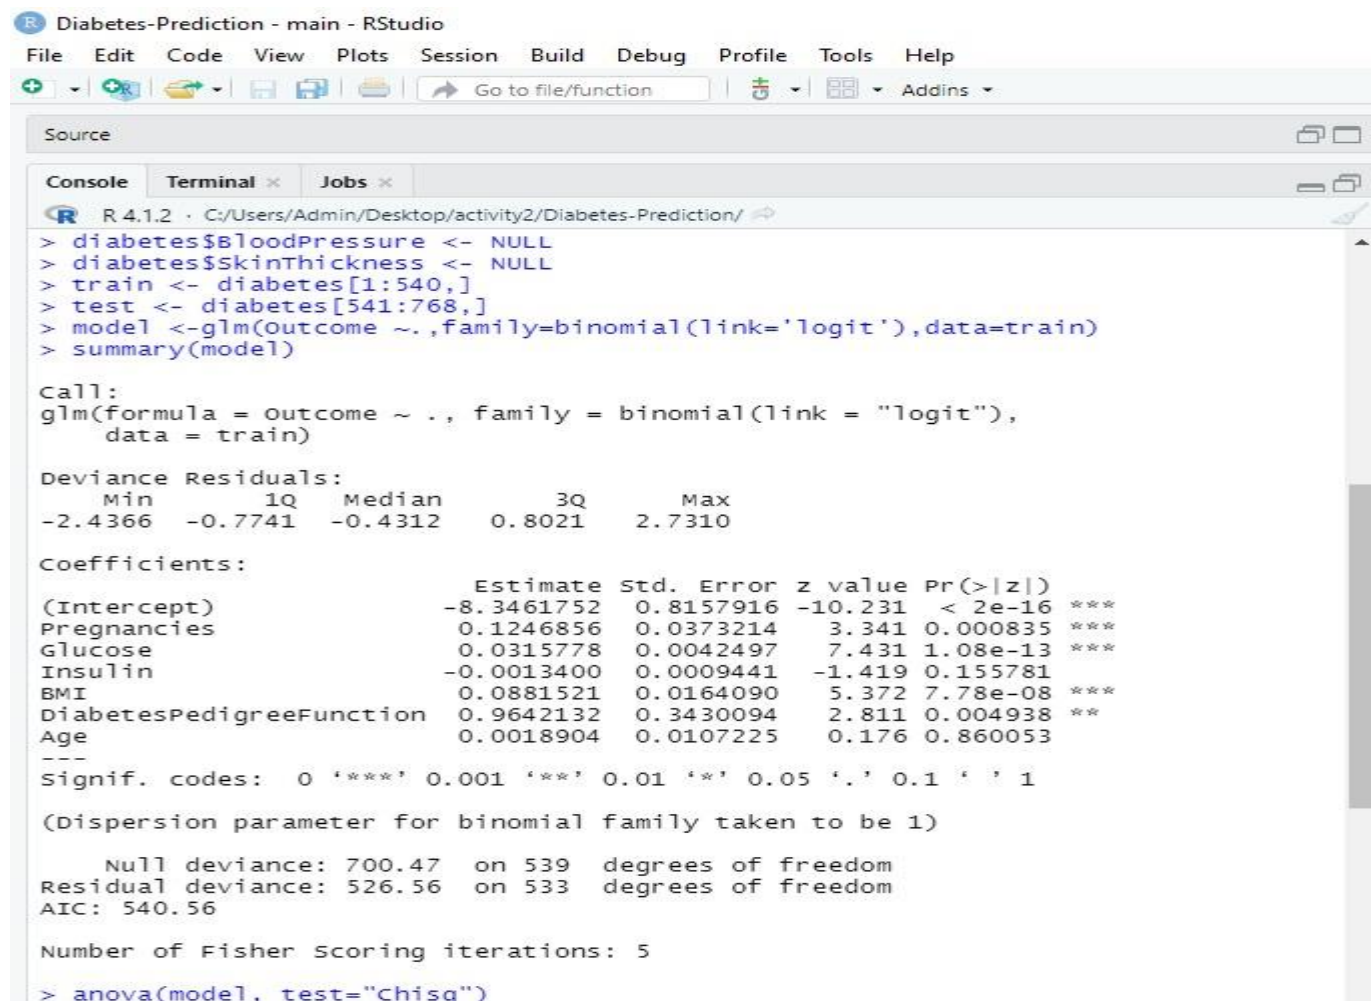
model <- glm(train ~ ., data = x,family='binomial')

summary(logistic)

#Predict Output

predicted= predict(logistic,test)
```

Output:



```
R 4.1.2 · C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/
> diabetes$BloodPressure <- NULL
> diabetes$SkinThickness <- NULL
> train <- diabetes[1:540,]
> test <- diabetes[541:768,]
> model <- glm(outcome ~.,family=binomial(link='logit'),data=train)
> summary(model)

Call:
glm(formula = outcome ~ ., family = binomial(link = "logit"),
    data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.4366  -0.7741  -0.4312   0.8021   2.7310

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)   -8.3461752   0.8157916  -10.231  < 2e-16 ***
Pregnancies     0.1246856   0.0373214   3.341 0.000835 ***
Glucose         0.0315778   0.0042497   7.431 1.08e-13 ***
Insulin        -0.0013400   0.0009441  -1.419 0.155781
BMI             0.0881521   0.0164090   5.372 7.78e-08 ***
DiabetesPedigreeFunction 0.9642132   0.3430094   2.811 0.004938 **
Age            0.0018904   0.0107225   0.176 0.860053
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 700.47  on 539  degrees of freedom
Residual deviance: 526.56  on 533  degrees of freedom
AIC: 540.56

Number of Fisher Scoring iterations: 5
> anova(model, test="Chisq")
```


The result shows that the variables Triceps_Skin, Serum_Insulin and Age are not statistically significant. In other words, the p-values are greater than 0.01. Therefore they will be removed.

Diabetes-Prediction - main - RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

Source

Console Terminal x Jobs x

R 4.1.2 · C:/Users/Admin/Desktop/activity2/Diabetes-Prediction/

```
> anova(model, test="Chisq")
Analysis of Deviance Table

Model: binomial, link: logit
Response: Outcome
Terms added sequentially (first to last)
```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)	
NULL			539	700.47		
Pregnancies	1	26.314	538	674.16	2.901e-07	***
Glucose	1	102.960	537	571.20	< 2.2e-16	***
Insulin	1	0.062	536	571.14	0.803341	
BMI	1	36.135	535	535.00	1.841e-09	***
DiabetesPedigreeFunction	1	8.414	534	526.59	0.003723	**
Age	1	0.031	533	526.56	0.860201	

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> fitted.results <- predict(model,newdata=test,type='response')
> fitted.results <- ifelse(fitted.results > 0.5,1,0)
> misClasificError <- mean(fitted.results != test$Outcome)
> print(paste('Accuracy',1-misClasificError))
[1] "Accuracy 0.789473684210526"
```

Accuracy of Logistic Regression is 78.94%

Decision Tree:

Now we present the Decision Trees algorithm

Algorithm of Decision Tree In R:

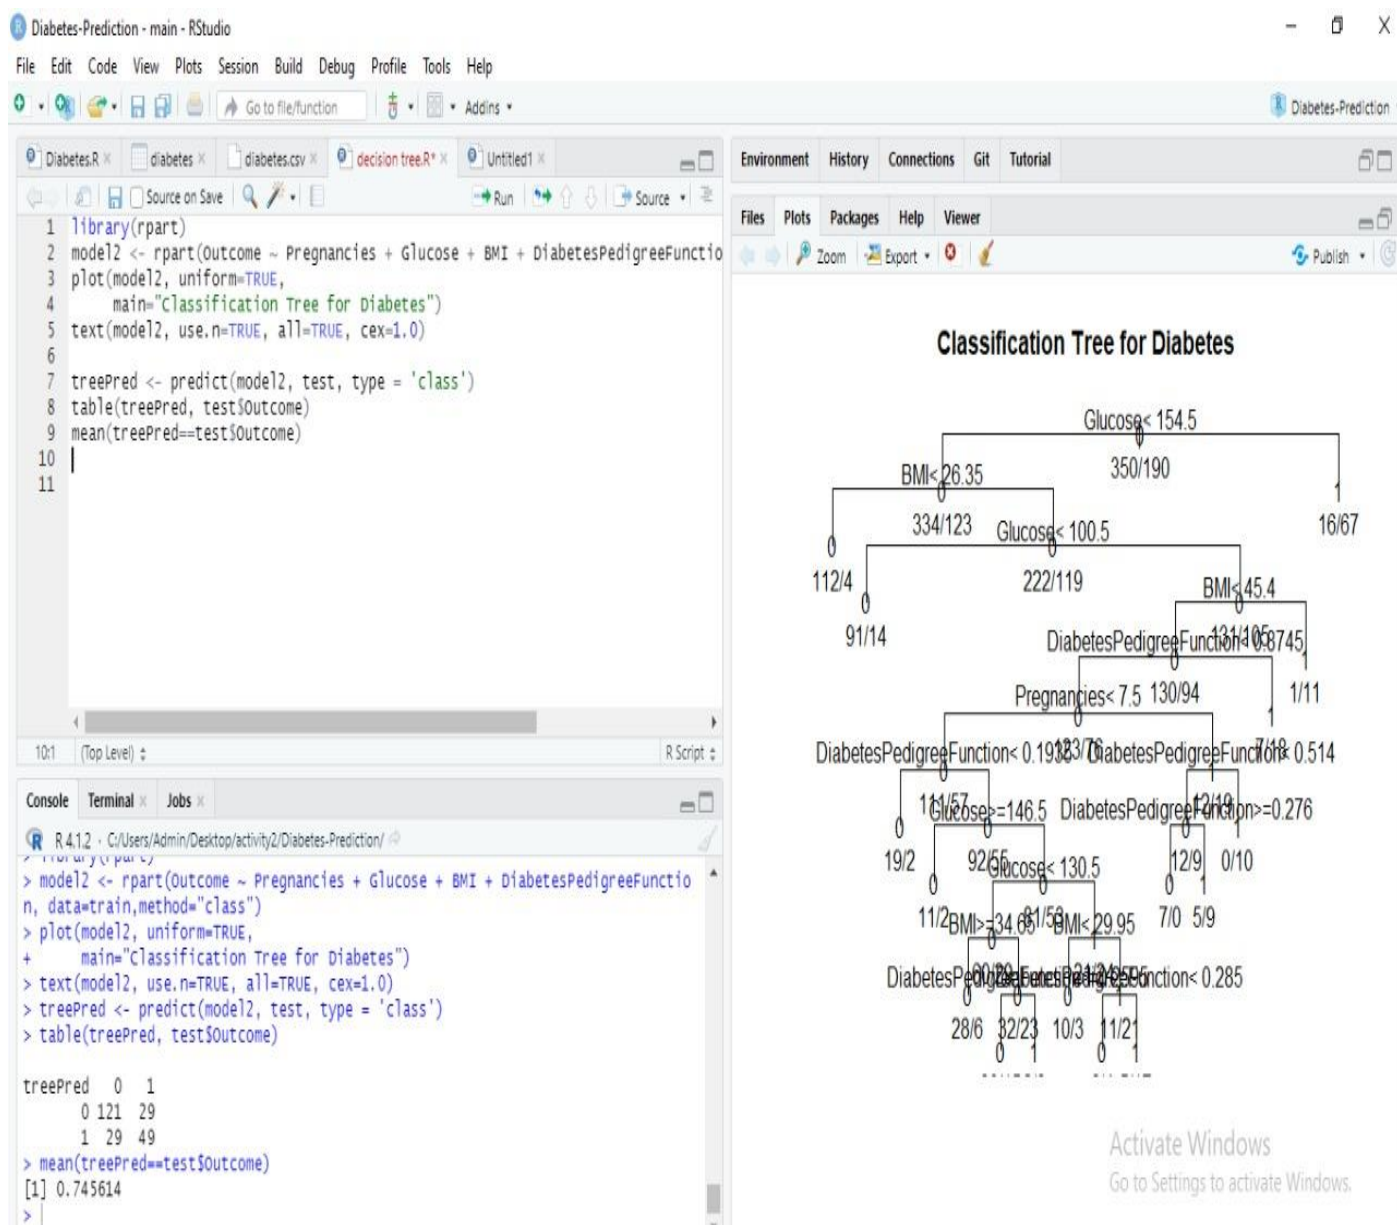
```
library(rpart)
x <- cbind(x_train,y_train)
# grow tree
fit <- rpart(y_train ~ ., data = x,method="class")
summary(fit)
#Predict Output
predicted= predict(fit,x_test)
```

Output:

The results display the split criterion (e.g. Plasma_Glucose < 154.5), the number of observations in that branch, the deviance, the overall prediction for the branch (Yes or No), and the fraction of observations in that branch that take on values of Yes and No. Branches that lead to terminal nodes are indicated using asterisks.

Now we plot of the tree, and interpret the results.

“Diabetes” appears to be Plasma_Glucose, since the first branch split criterion (e.g. Plasma_Glucose <154.5).



The test error rate is 30%. In other words, the accuracy is 70%.